

# Comparative analysis of Gene Finding tools when applied to *Trichoderma* genomes

Connor Burbridge<sup>1</sup>

<sup>1</sup>USask NSID: cbe453 , USask ID no. 11162928 , Supervisors:  
Dave Schneider & Tony Kusalik,

February 15, 2024

# 1 Permission stuff

## 2 Abstract

### 3 Acknowledgements

# Contents

<b>1</b>	<b>Permission stuff</b>	<b>2</b>
<b>2</b>	<b>Abstract</b>	<b>3</b>
<b>3</b>	<b>Acknowledgements</b>	<b>4</b>
<b>4</b>	<b>Background</b>	<b>7</b>
4.1	Trichoderma . . . . .	7
4.2	Genome Assembly . . . . .	7
4.3	Repeat Identification/Masking and Identification of AT-rich Genomic regions . . . . .	8
4.4	Gene Finding . . . . .	9
<b>5</b>	<b>Platform and Software Installation</b>	<b>10</b>
5.1	Platform . . . . .	10
5.2	NextDenovo and NextPolish Installation . . . . .	10
5.3	RepeatMasker Installation . . . . .	10
5.4	GeneMark-ES Installation . . . . .	11
5.5	Braker2 Installation . . . . .	11
<b>6</b>	<b>Methodology</b>	<b>11</b>
6.1	Methodology Overview . . . . .	11
6.2	Selection of Existing <i>Trichoderma</i> Genome Assemblies . . . . .	11
6.3	Assembly . . . . .	13
6.4	Repeat Masking . . . . .	13
6.5	GeneMark-ES . . . . .	14
6.6	Braker2 . . . . .	14
<b>7</b>	<b>Identification of Overlapping Features and Regions</b>	<b>15</b>
<b>8</b>	<b>Analysis of Results</b>	<b>15</b>
8.1	Basic Analysis . . . . .	16
8.2	Distribution of Gene Lengths . . . . .	16
8.3	Intersection of Gene Calls, smallRNAs and Repetitive Regions	17
8.4	Methodology for Indetifying Overlapping Features . . . . .	18
8.5	Shared Gene Content with Yeast . . . . .	18
8.6	Comparative Genomics . . . . .	19

<b>9</b>	<b>Results</b>	<b>19</b>
9.1	Assemblies of DC1 and Tsth20 . . . . .	19
9.2	Initial Gene Finding Results . . . . .	19
9.3	Region Identification . . . . .	20
<b>10</b>	<b>Discussion</b>	<b>22</b>
<b>11</b>	<b>Conclusions</b>	<b>22</b>
<b>12</b>	<b>References</b>	<b>24</b>

## 4 Background

### 4.1 Trichoderma

Crop resistance to environmental stressors is a necessity for crop health and overall crop yields. Current popular methods for crop protection involve the use of pesticides and genetically modified organisms, which can be expensive and potentially politically dividing in the case of GMOs[1]. In addition, crops suffer when soils are not sufficient for crop growth and health. Soil insufficiencies can result in drought stress as well as nutrient stress, leading to poor overall yields.

*Trichoderma* is a type of fungi that can colonize the roots of plants in a non-toxic, non-lethal, opportunistic symbiotic relationship[2]. Many strains of *Trichoderma* have been shown to provide resistance to bacteria and other fungi in soils through the use of polyketides, non-ribosomal peptide synthetases and other antibiotic products[2][3]. Recently, two strains of *Trichoderma* have been identified in the prairie regions of Alberta and Saskatchewan. These two strains, named Tsth20 and DC1, have been found to have beneficial properties when used as an inoculant for plants in the soils mentioned before. In addition to these beneficial properties, the two strains mentioned previously provide even further protection for plants in dry, salty soils and one strain also has potential for use as a bioremediation tool in soils contaminated with hydrocarbon content. Bioremediation and resistance to drought tolerance has also been investigated in other strains of *Trichoderma* as well[4][5]. However, little is known about the mechanisms at work in these strains, so DC1 and Tsth20 were sequenced by the Global Institute for Food Security (no publication yet) in an initial attempt to better understand the details of these genomes. While this research does not directly identify genomic elements related to the secretome of these genomes, it may serve as a foundation for future research of *Trichoderma*.

### 4.2 Genome Assembly

Sequence assembly has been a long-standing application problem in the field of bioinformatics[12]. Determining the correct order and combination of smaller subsequences into an accurate complete sequence assembly is computationally difficult in terms of compute resources such as memory, CPU cycles and storage required for input sequences[12]. In addition to these difficulties,

there can be other issues encountered during assembly due to the nature of the data or genomes themselves, such as low quality base calls for long read data or the inherent content of genomes themselves using repetitive regions as an example. Insufficient data used in an assembly may result in short, fragmented assemblies, depending on the size of the genomes, while sequence data that is not long enough can fail to fully capture repetitive regions in an assembly. To solve this problem, a wide range of assembly tools have been developed with their own unique approaches to the genome assembly problem, so it is important to use an appropriate assembler for the task at hand, and also important to evaluate the assembly thoroughly. One approach to aid in the previously mentioned issue of assembly correctness is to use a combination of long and short reads in what is known as a hybrid assembly. Combining both highly accurate short reads with deep coverage along with less accurate but much longer reads can produce high quality genome assemblies that capture long repetitive regions. Hybrid assembly approaches have been shown to produce high quality assemblies in a wide variety of organisms as they combine long read data with short data to produce assemblies that properly represent long repetitive regions with additional high quality Illumina sequences for correction. Once assembled, the sequences must also be evaluated with measures such as N50, L50, coverage, average contig length and total assembled length to ensure that the genomes are well assembled, at least based on these metrics[12]. Following appropriate assembly protocols is essential to the further success of a project as downstream processing such as annotation depends on a high-quality assembly.

### **4.3 Repeat Identification/Masking and Identification of AT-rich Genomic regions**

Repeat identification within assembled genomes is a problem that needs to be considered during the genome annotation process. Regions with long repeats can have a significant impact on genome assembly as well as gene finding due to the limitation of short reads used in some assemblies[14]. Short reads may be unable to bridge or cover entire repeat regions within a genome, so it is important to consider the use of long reads from technologies such as Nanopore or PacBio to provide a complete picture of these regions when pursuing a new genome assembly project. It is also possible for repetitive regions to contain genes as well, making for an interesting investigation in



regards to *Trichoderma*, as fungal genomes have been shown to contain many repeat regions with a high concentration of A and T nucleotides[15]. Once these repetitive regions have been identified, the genome could be masked to exclude these regions in downstream processing if desired, as these regions may be poorly assembled and may result in found genes that do not truly exist in those regions. However, this may not be as common today, as repetitive regions have been shown to contain genes as well[16]. This may affect the gene finding process described later and may be an interesting topic to look into considering the large number of available gene finding programs.

## 4.4 Gene Finding

Gene finding is another long standing computational problem in bioinformatics, which concerns itself with identifying potential genes within genomes based on patterns or evidence considered by the gene finding program. There are two common methods for gene finding, those methods being *ab initio* methods, where programs search for patterns and gene structures, and similarity or evidence-based searches, which use prior information such as RNAseq data, expressed sequence tags and protein sequences to identify genes within a new genome[17]. Complicating the process more is the introduction of introns and alternative splicing in eukaryotes, making it possible for one gene to have several possible transcripts at the same locus. Examples of *ab initio* methods include tools such as GeneMark-ES[7] and GlimmerHMM[8], while evidence based methods include tools such as Braker2[6], which can incorporate existing data such as RNAseq, protein sequences, etc. in gene prediction models.

As mentioned previously, long repeats and transposable elements can cause issues for gene finding programs, making this an interesting area of research, at least for fungal genomes such as *Trichoderma*. Fungal genomes are also interesting in the topic of gene finding as there are few programs targeted directly at gene finding in fungal genomes while organisms such as human, mouse, and *Arabidopsis* benefit from having many tools tested on them as they are model organisms in the field, according to table 1 of Wang, Z *et al*[17]. How these different methods and tools perform when applied to fungal genomes is an important consideration as fungi have features that can benefit plant growth as mentioned earlier.

There are also other aspects of gene finding tools that are important to consider. These include features such as whether or not the gene finders find

non-coding RNAs, annotation of 5' and 3' UTR regions, and in the case of ab-initio methods, the assumptions made by the underlying models used for gene finding. These features and others can influence a user's decision on which gene finding tool to consider and will complicate comparative analysis of multiple gene finding tools.

## 5 Platform and Software Installation

(Possibly supporting materials or discussion)

### 5.1 Platform

All analysis was performed on the RSMI server hosted on Copercius at the University of Saskatchewan. This server is equipped with 64 cores in addition to 1.5 TB of memory. The server is running RedHat Enterprise Linux 7 as of writing this thesis. All data is stored either on datastore, or in the RSMI scratch space.

### 5.2 NextDenovo and NextPolish Installation

Installation of nextDenovo was straightforward. Simply download the compressed tar file from their website and unpack it. NextDenovo requires Python versions 2 and 3 along with a package called parallel to aid in parallel processing of datasets. The parallel package was installed using pip in the bioinformatics conda environment in the scratch space of Copernicus. NextPolish was installed in a Python environment by a member of the research computing team that manages our system. Assistance was required for this as the version of RHEL used by the server introduces glibc version conflicts with Anaconda when trying to install nextPolish.

### 5.3 RepeatMasker Installation

The installation procedure was somewhat indepth, requiring RepeatMasker configuration, which itself requires downloading an appropriate repeat database (Dfam in this case, included with RepeatMasker), installation of Tandem Repeat Finder (TRFM) and installation of a sequence search tool, for which I chose HMMER from the list of potential tools as we were generally familiar

with its use. The path to the installation of TRFM is required during configuration along with the search tool of choice, a simple selection of 4 tools that will have an autocompleted path in this case, since HMMER is installed via anaconda.

## 5.4 GeneMark-ES Installation

GeneMark-ES was successfully installed by downloading and unpacking the package from their website along with a key required for use.

## 5.5 Braker2 Installation

Braker2 was also successfully installed by a member of the research computing team who has set up several modules including an initialization script to get things up and running as well as create a reloadable environment for use again in the future. Once the environment has been loaded, one must load the Hisat2 module from Compute Canada as well as an htlib module (more detail to come). Once all modules are loaded, there are a few environment variables that need to be set, those being AUGUSTUS\_CONFIG\_PATH and TSEBRA\_CONFIG(?figure this out). In addition, a software package named TSEBRA from the same developers as Braker2 must be installed for consolidating gene calls. The variables can be set within the braker2.pl command, which have higher priority over environment variables and probably makes things easier to track.

# 6 Methodology

## 6.1 Methodology Overview

The general methodology for this work is described in figure 1. Each portion of this figure is discussed in detail in this section.

## 6.2 Selection of Existing *Trichoderma* Genome Assemblies

Comparing the performance and features, of gene finding tools, both qualitative and quantitative, in the context of any set of genomes is important




Figure 1: A flowchart of the methodology followed for this research. Sections are separated based the general process they are associated with (i.e. input data, assembly, gene finding and downstream analysis).

for those interested in selecting a specific gene finding tool. To accent(?) the processing for genomes of interest, those being DC1 and Tsht20, we should include other previously assembled *Trichoderma* assemblies. Currently selected genomes include *Trichoderma reesei*, *Trichoderma harzianum*, and *Trichoderma virens*, with *Trichoderma reesei* being the 'reference' in this case, as it is well studied and there are several patents involving it's use as an organism for production of compounds such as antibiotics in industrial applications.

### 6.3 Assembly

In an attempt to produce high quality assemblies of DC1 and Tsth20, We decided on a set of tools named NextDenovo and NextPolish as they have produced excellent assemblies based on previous experience. (should find a citation to confirm this)

(Might be better for discussion or omitted since it is specific to our setup) Initial attempts to run the example dataset resulted in permissions errors due to the management of the storage system being used, which were encountered with other tools in the past. To remedy this, the software installation was copied to RSMI's scratch space on Copernicus. Once the appropriate permissions were given to run nextDenovo, the example dataset was run without issue.

Following assembly using nextDenovo, Illumina sequence data from DC1 and Tsth20 was used to polish each respective genome using nextPolish. Default parameters were used from assembly except for modification of the parallel option to reduce processing times.

### 6.4 Repeat Masking

In order to evaluate the performance of gene finding tools in repetitive or low complexity regions in the context of *Trichoderma* genomes, we must first identify said regions in the genomes considered. To do this, RepeatMasker has been selected as a tool to identify repeat regions based on a fungal subset of the Dfam database by specifying the fungi species tag to RepeatMasker when running the program. The program was configured with options to produce several output formats for each genome considered, which will allow for more informative downstream analysis of results. All commands for repeat masking are located within the processing directory for each strain/genome.

(probably more suited for additional materials at the end?) General command for running RepeatMasker:

```
/datastore/Roots/Connor/masters/software/repeatmasker/RepeatMasker/RepeatMasker  
-pa 10 -a -small -species fungi -html -gff -dir ./ path-to-genome/genome.fasta
```

## 6.5 GeneMark-ES

To begin, GeneMark-ES was run as it requires no prior information or alignments in order to run. In this case GeneMark-ES has an option specifically for fungal genomes, which was used in this case. Apart from the fungal option, the only additional options supplied were for output format of GFF3 and number of cores for reduced processing time.

General command structure for GeneMark-ES:

```
gmes_petap.pl -ES -fungus -format gff3 -cores 48 -sequence /path/to/sequence
```

## 6.6 Braker2

As mentioned previously, *Trichoderma reesei* was selected as the reference genome for this work. With this in mind, several short read archives (SRAs) from *T. reesei* were selected for Augustus training. Following Augustus training, the model for *T. reesei* was applied to all genomes considered. Settings and procedures from running Braker2 are described below.

The variables that need to be set are AUGUSTUS\_CONFIG\_PATH and TSEBRA\_PATH. Augustus, by default, tries to write species information to the location where the software is installed. In this case, we don't have write permissions to the compute canada software stack hosted by Research Computing, so the AUGUSTUS\_CONFIG\_PATH variable must be set in order to create a writeable directory. As long as that path has a directory within it called braker, and a species directory within the braker directory, things should go smoothly. TSEBRA is a set of scripts also made by the creators of Braker and is required to merge results from the various gene prediction tools involved in the Braker2 pipeline. The TSEBRA\_PATH simply points to the directory where TSEBRA is located Both Braker2 and TSEBRA can be cloned directly from GitHub (links to come)

```
iiiiii HEAD
```

## 7 Identification of Overlapping Features and Regions

**Feature Identification:** To first understand how gene prediction tools perform in comparison to other gene prediction tools, we must identify features. This identification of features will help us describe the similarities, and differences between gene finding tools. A feature, in this context, is any feature stated within a Genomic Feature Format file (GFF) provided to the program, in which multiple GFF files can be provided. The definition of a feature, for this application, is an object that contains a contig ID, a start position, an end position and a strand property. In the context of features on different strands, start and stop positions of features are sorted based on left and right positions of the feature in respect to the reference sequence.

**Region Identification:** In addition to feature creation, we will also identify regions of overlapping features based on the predictions from each gene finding tool. These regions will help identify the agreements, or disagreements, between different gene-finding tools. A region, in this context, is a set of overlapping features, all of which overlap at least one other feature in the region. With each overlap, there will be an overlap type. These types can be defined based on Allen's Interval Calculus (reference), with the exception of features that start beyond the end point of the current region.

===== Example command for braker2:

```
/scratch/p2irc/p2irc_rsmi/cbe453/masters/software/braker2/BRAKER/scripts/braker.pl  
-gff3 -threads 60 -TSEBRA_PATH=/scratch/p2irc/p2irc_rsmi/cbe453/masters/software/braker2/  
-genome /path/to/sequence -species=TriseiFungal -fungus -useexisting
```

## 8 Analysis of Results

After completion of the processing portion of this work, the results must be processed in a useful way, which includes both the biological implications of the gene calls as well as the computational, or gene finding features, of the the selected programs. To better understand how gene finders perform in these two classes, we must define an appropriate plan for analysis of the results produced so far. Currently, downstream analysis plan has been broken down into several sections.

## 8.1 Basic Analysis

Basic analysis of gene finding results is an important part of this research. Total gene, transcript and protein counts will be identified for each genome and gene finding tool combination. Comparing the general outputs of these programs will provide an idea of their performance in different *Trichoderma* genomes. In addition to these basic outputs, analysis will also be performed for the following: distribution of gene lengths, intersection of gene calls, smallRNAs and repetitive regions, shared gene content with a close fungal relative. Analysis for these results can be performed through simple shell scripting with grep and other unix tools, although processing through Python might provide results that are easier to reproduce with proper programming. Having one script with several modules that can be rerun at will would be easier to handle than multiple shell scripts. This thinking for processing will be applied to subsequent sections of this as well.

## 8.2 Distribution of Gene Lengths

One important aspect of gene finding tools to consider is the distribution of gene lengths predicted by each individual tool. Certain tools, such as GeneMark are based on pre-defined models, which may limit the length of predicted genes, while tools such as Braker2, which incorporate RNAseq data, may predict a wider distribution of gene lengths depending on the input dataset used. Regardless, the ability of a gene finding tool to predict a wider range of gene lengths can be useful if users are looking for short or larger genes. To help determine whether or not these tools find shorter genes, or small RNAs, the genomes of interest have been annotated using Infernal along with the Rfam database to identify small RNAs as a ground truth. These annotation results will also be included with results from other annotation processes further down the line. Again, these results can be produced with a Python script. The resulting data could then be used as input to violin plots for each genome and set of tools considered in this analysis process. Violin plots should provide a good visualization of gene lengths as well as the number of genes found with specific lengths. Means could also be compared statistically for genomes and the multiple tools considered as well.

Analysis of gene lengths was performed using a Python script. Combined predicted CDS sequences for each predicted gene were used as input for the



total gene length. CDS sequences predicted by Braker2, were directly available in the output directories when the program was run. CDS sequences from GeneMark required extraction of the CDS sequences from the genome FASTA files. This process was performed using the gffread tool from the Cufflinks package. Predicted CDS sequences were loaded into Python using Biopython's SeqIO package. Sequence lengths were then placed in a list and analyzed using a combination of pandas and numpy. A log10 transformation was applied to the sequence lengths as the original distribution was heavily skewed due to long outlier CDS sequences. After transformation, the CDS length distribution appears as a normal distribution, although there are interesting troughs that occur in several of the peaks for several of the genomes considered. These troughs did not appear in a comparable Yeast reference dataset, although the log transformed data still appears to be a normal distribution.

### 8.3 Intersection of Gene Calls, smallRNAs and Repetitive Regions

Annotation of all three features in the title are important in assessing the ability of gene finding tools. Even more important, is the potential for overlap between gene calls and small RNAs as well as repetitive regions the genomes. As discussed in the previous subsection, the distribution of gene lengths predicted by a gene finding tool can be an important metric for users. Overlapping predicted genes from tools alongside the output from Infernal and the Rfam database may provide insight into whether or not these gene finding tools are able to predict RNAs of very short length. In addition to small RNAs, repetitive regions in *Trichoderma* genomes hold potential for recombination and gene content, although the inherent nature of these repetitive regions (low nucleotide diversity) suggests that gene content should be low, based on the nucleotides required for start and stop codons. Analysis of these intersections can be performed via bedtools or through biopython (I believe). Again, having all processing steps included in one script as separate functions that can be called at whim will make further processing easier if changes need to be made.

## 8.4 Methodology for Identifying Overlapping Features

To analyze the results from multiple gene-finding tools, we must first define two conceptual topics.

Definition of a Feature: First, a feature, in the context of this research, is any item contained within a Genomic Feature Formatted file (GFF). Each feature contains a contig ID, a start position, end position, and a feature ID based on the information from the GFF file. These features will be used in the process of identifying regions, or overlapping features from prediction tools.

Definition of a Region: Secondly, a region is defined as any overlap between features relative to the reference sequence being considered. To identify regions, every feature from each GFF file being considered, is sorted by start position. For clarification, the start position is based on the left most position in the GFF file, regardless of the strand that the feature is predicted on. Once the features have been sorted by left position, the sorted features are iterated over to identify regions, as long as the left position of the next feature is consistent with the left position, within the left and right position, or equal to the right position of the current region. One drawback to this approach is that ideally, identification of overlap types based on Allen's interval algebra would be performed at this point. However, the methodology of initially sorting features based on left position somewhat prevents this process from happening. This implementation requires further processing of identified regions late on the process, simplifying the all to all comparison that would occur if all features were considered at once.

## 8.5 Shared Gene Content with Yeast

While considering novel gene calls can be useful, comparing those calls to a well-studied close relative can provide a rudimentary validation of the calls as a ground truth. In this case, a comparison to Yeast will be made. The agreed upon number for successful gene-finding as compared to Yeast is roughly 80-85% of gene content. This will confirm that at least most of a closely related fungal genome's content is predicted and shared by the gene calls for *Trichoderma*. Results for this processing can be produced with a simple BLAST search and appropriate cutoff values (i.e. query coverage, percent nucleotide identity, E-score, etc.). Other tools for evaluation will certainly be considered, although I will need to look into this process further.

## 8.6 Comparative Genomics

With the data produced by this research, it is possible to perform some comparative genomics (time permitted), mostly related to the assemblies generated during this work along with the RefSeq genomes included from NCBI. Mummer is a potential tool to use for all to all genome alignments, although there may be difficulty in the ordering of contigs/scaffolds/chromosomes when performing these alignments. This work is not necessarily required but would be interesting from a biological perspective to identify rearrangements, inversions and such.

## 9 Results

### 9.1 Assemblies of DC1 and Tsth20

From figure 2, we can make several observations, the first of which is total contig counts for each assembly. For DC1 and Tsth20, the total contig counts are an order of magnitude smaller when compared to the other NCBI RefSeq assemblies, indicating highly contiguous assemblies from nextDenovo and nextPolish. This is likely due to the use of long-read sequencing used in the assemblies of DC1 and Tsth20. For all assemblies, the total assembly lengths are similar, hovering around the 38-42Mb range, except for *T. reesei*, which is known to have a significantly smaller genome length (ref) at roughly 33Mb. The largest contig size for each assembly vary greatly. DC1 and Tsth20 have the largest contigs of all assemblies being considered, which is again likely due to the inclusion of long-read sequencing data in the assembly process. The N50 values for all assemblies are above 1Mb, with DC1 and Tsth20 N50s being at minimum three times larger than others assemblies.

### 9.2 Initial Gene Finding Results

In figure 3, we see a summary of total genes and mRNAs predicted by each of the selected tools. An immediate trend can be seen in this data. The total predicted features for Braker2 are significantly lower than those from GeneMark in all assemblies except for *T. reesei*. This may be due to Braker2 using RNaseq data from *T. reesei* during its training process, which we will cover in the discussion section. Another observation can be seen when comparing the predicted genes and predicted RNAs for the same tool when applied to the

Strain	Total Contigs	Total Length	Largest Contig	GC%	N50	L50
DC1	8	38.6 Mb	11.49 Mb	47.97	5.69 Mb	3
Tsth20	7	41.58 Mb	8.02 Mb	47.33	6.52 Mb	3
<i>T. harzianum</i>	532	40.98 Mb	4.08 Mb	47.61	2.41 Mb	7
<i>T. virens</i>	93	39.02 Mb	3.45 Mb	49.25	1.83 Mb	8
<i>T. reesei</i>	77	33.39 Mb	3.75 Mb	52.82	1.21 Mb	9

Figure 2: General assembly metrics produced by QUAST (a genome quality assement tool).

assemblies. GeneMark does not appear to identify any additional isoforms, only reporting the entire gene structure. Braker2 does identify isoforms, although very few of them. This may be related to the training set provided to the training process, although this is not yet confirmed. Overall, it appears that GeneMark regularly predicts a higher number of features when applied to these assemblies.

One important aspect to consider when looking at the output of different gene prediction tools is the distribution of sequence lengths predicted by any given tool. Lengths of possible sequences can vary widely, ranging from small non-coding RNAs, which can be less than 200 nucleotides in length, up to the largest genes which cover more than two kilobases. Due to this wide variation in possible sequence lengths, it is possible that different prediction tools could produce different distributions of predicted sequence lengths. This is important if researchers are interested in small non-coding RNAs or atypically large genes. This section will investigate the coding sequence lengths predicted by Braker and GeneMark for DC1 and Tsth20 while the RefSeq assemblies will also include the RefSeq annotation.

Figures here

### 9.3 Region Identification

Based on the methodology described earlier, we have identified sections of contigs with one or more overlapping predicted sequences (features). The result of this processing is a GFF formatted file with each entry containing information about a region on a specific contig. Also contained in each entry, is a list of all features included in that region with start/stop position and




Figure 3: Shows the counts of genes and mRNAs found by each gene finding tool for each genome assembly considered.

Assembly	Regions (total)	RefSeq	Braker2	GeneMark
DC1	11268	NA	8485	11265
Tsth20	12285	NA	8747	12279
<i>T. harzianum</i>	13377	13122	8051	11529
<i>T. virens</i>	12045	11864	7554	11413
<i>T. reesei</i>	9818	8836	9202	8897

Figure 4: Counts of regions identified in total and total number of regions where a prediction from each individual tool was found.

strand, the tools used to predict the features, and a count of how many prediction tools that were included in said region. Initial analysis of the output is discussed in this section.


With our results, the first interesting area to look at is simply the total number of regions identified for any given assembly as well as number of regions that each tool was a member(weird wording) of. In figure 4, we see that Braker2 deviates from both the RefSeq and GeneMark predictions greatly, except in the case of *T.reesei*. As speculated earlier, this may be due to the training dataset supplied to Braker2 as it happens to be present in a similar number of regions when compared to RefSeq and GeneMark in *T. reesei*. In general, it appears that the RefSeq predictions are present in more regions than both Braker2 and GeneMark.

Figure 4 demonstrates that the composition of regions can differ significantly based on the tools used. To illustrate agreement of tools in regards to identified regions, Venn diagrams were generated showing agreement of predictions in a region for each tool.


blah blah

## 10 Discussion


## 11 Conclusions



(a)



(b)



(c)

## 12 References

### References

- [1] Zhang, C., Wohlhueter, R., Zhang, H. (2016). Genetically modified foods: A critical review of their promise and problems. *Food Science and Human Wellness*, 5(3), 116-123.
- [2] Hermosa, R., Viterbo, A., Chet, I., Monte, E. (2012). Plant-beneficial effects of *Trichoderma* and of its genes. *Microbiology*. 38, 17-25.
- [3] Ramirez-Valdespino, C., Casas-Flores, S., Olmedo-Monfil, V. (2019). *Trichoderma* as a Model to Study Effector-Like Molecules. *Frontiers in Microbiology*, 15.
- [4] Repas, T.S., Gillis, D.M., Boubakir, Z., Bao, Xioahui., Samuels, G.J., Kaminskyj, S.G.W. (2017). Growing plants on oily, nutrient-poor soil using a native symbiotic fungus. *PLOS ONE*, 12(10).
- [5] Kaminskyj et al. (2012). Method for increasing plant growth using the fungus *Trichoderma harzianum* (United States PCT/CA2010/OO1454). <https://patentimages.storage.googleapis.com/57/4f/7f/5ff60c47045f01/US20120178624A1.pdf>
- [6] Hoff, K. J., Lomsadze, A., Borodovsky, M., Stanke, M. (2019). Whole-Genome Annotation with BRAKER. *Methods Mol Biol.*, 1962, 65-95.
- [7] Borodovsky, M., Lomsadze, A. (2011). Eukaryotic Gene Prediction using GeneMark.hmm-E and GeneMark-ES. *Current Protocols in Bioinformatics*. 4.
- [8] Majoros, W.H., Pertea, M., Salzberg, S.L. (2004) TigrScan and GlimmerHMM: two open-source *ab initio* eukaryotic gene-finders. *Bioinformatics*. 20(9), 2878-2879.
- [9] Stanke, M., Morgenstern, B. (2005). AUGUSTUS: a web server for gene prediction in eukaryotes that allows user-defined constraints. *Nucleic Acids Research*. 33, 465-467.



- [10] Jones, P., Binns, D., Chang, H. Y., Fraser, W., Li, W., Hunter, S. (2014). InterProScan 5: genome-scale protein function classification. *Bioinformatics*. 30(9), 1236-1240.
- [11] Yandell, M., Ence, D. (2012). A beginner's guide to eukaryotic genome annotation. *Nature Reviews Genetics*. 13, 329-342.
- [12] Sohn, J., Nam, J. (2016) The present and future of *de novo* whole-genome assembly. *Briefings in Bioinformatics*. 191, 23-40.
- [13] Bankevich, A., Nurk, S., Antipov, D., et al. (2012). SPAdes: a new genome assembly algorithm and its applications to single-cell sequencing. *Journal of Computational Biology*. 19(5), 455-477.
- [14] Lerat., E. (2009). Identifying repeats and transposable elements in sequences genomes: how to find your way through the dense forest of programs. *Nature Heredity*. 104, 520-533.
- [15] Li, WC., Huang, CH., Chen, CL. et al. (2017). *Trichoderma reesei* complete genome sequence, repeat-induced point mutation, and partitioning of CAZyme gene clusters. *Biotechnol Biofuels*. 10, 170.
- [16] Slotkin, R.K. (2018). The case for not masking away repetitive DNA. *Mobile DNA*. 9(1), 15.
- [17] Wang, Z., Chen, Yazhu., Li, Y. (2016). A Brief Review of Computational Gene Prediction Methods. *Genomics Proteomics Bioinformatics*. 24, 216-221.
- [18] Smit, A., Hubley, R., Green, P. (2013-2015) RepeatMasker Open-4.0. <https://www.repeatmasker.org>
- [19] Nawrocki, E.P., Eddy, S.R. (2013). Infernal 1.1: 100-fold faster RNA homology searches. *Bioinformatics*. 29, 2933-2935.
- [20] Gremme, G., Brendel, V., Sparks, M.E., Kurtz, S. (2005) Engineering a software tool for gene structure prediction in higher organisms. *Information and Software Technology*. 47, 965-978.
- [21] Adav, S.S., Sze, S.K. (2014). Chapter 8 - *TRichoderma* Secretome: An Overview. *Biotechnology and Biology of Trichoderma*. 103-104.

~~~~~ b715090053dedb063b1eacf591f238315062ce83